



Kubernetes kubectl and Deployment

.NET

*The kubectl CLI tool
allows for control of
Kubernetes clusters.*

[HTTPS://KUBERNETES.IO/DOCS/REFERENCE/KUBECTL/OVERVIEW/](https://kubernetes.io/docs/reference/kubectl/overview/)

Azure Kubernetes Service (AKS)

<https://kubernetes.io/docs/setup/production-environment/turnkey/azure/#azure-kubernetes-service-aks>

<https://github.com/Azure/aks-engine/blob/master/docs/tutorials/README.md>

<https://docs.microsoft.com/en-us/azure/aks/intro-kubernetes>

- The Azure Kubernetes Service (AKS) offers simple deployments for Kubernetes clusters.
- AKS makes it simple to deploy a managed Kubernetes cluster in Azure.
- AKS handles much of the complexity and operational overhead of managing Kubernetes.
- Azure handles critical tasks like health monitoring.
- The Kubernetes masters are managed by Azure. You only manage and maintain the agent nodes.
- AKS lets you integrate with Azure Active Directory and use Kubernetes role-based access controls.



Azure Kubernetes Service (AKS)

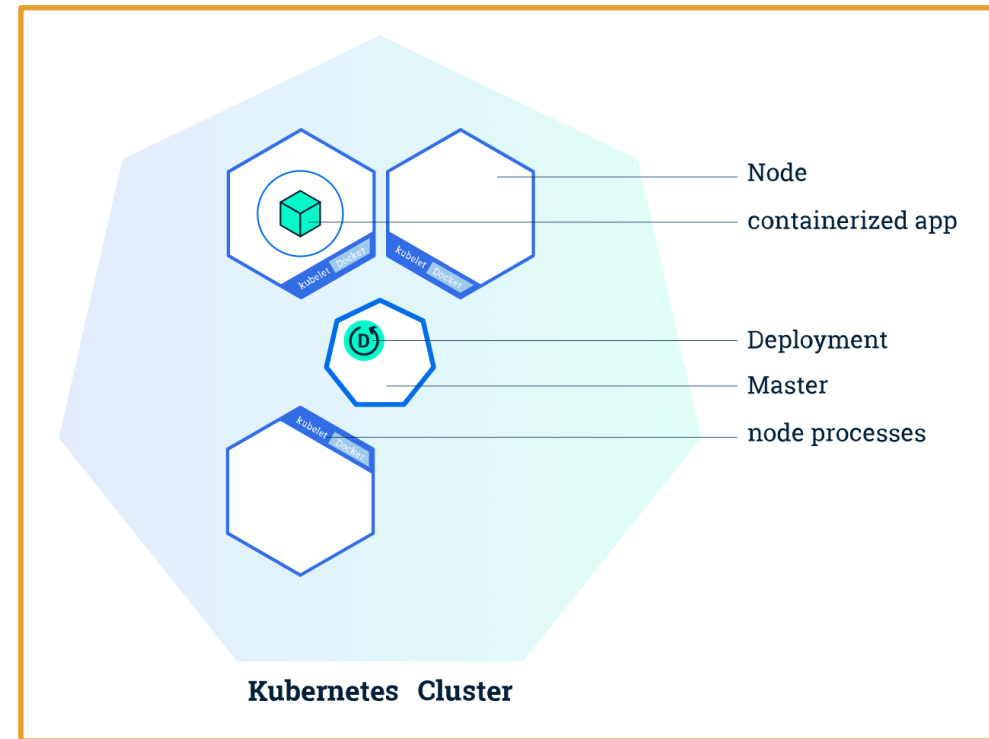
Kubectl (“Cube CTL”)

<https://kubernetes.io/docs/reference/kubectl/overview/>

Kubectl is the command line tool used to control Kubernetes clusters. **Kubectl** looks for a file named **config** in the **\$HOME/.kube** directory. You can specify other **kubeconfig** files by setting the **KUBECONFIG** environment variable or by setting the **--kubeconfig** flag.

Kubectl uses the **Kubernetes API** to interact with the cluster. The following syntax is used in command line to communicate through **kubectl**:

- **kubectl <command> <TYPE> <NAME> [flags]**



Kubectl

<https://kubernetes.io/docs/tasks/kubectl/install/>

<https://kubernetes.io/docs/reference/kubectl/cheatsheet/>

Kubectl uses the ***Kubernetes API*** to interact with the cluster. The following syntax is used in command line to communicate through ***kubectl***:

• **kubectl <command> <TYPE> <NAME> [flags]**

Command	Usage
command	Specifies the operation to perform on resources (create , get , describe , delete .)
type	Specifies the (case-insensitive) resource type.
--name	Specifies the case-sensitive name of the resource. If the name is omitted, details for all resources are displayed.
--flags, -f	Specifies optional flags.

Deployment (1 / 2)

<https://kubernetes.io/docs/tutorials/hello-minikube/>

A **Deployment** is the recommended way to manage the creation and scaling of **Pods**.

The **Deployment Controller**:

- uses a **Deployment .yaml** to change an unacceptable state to a desired state at a controlled rate.
- manages **ReplicaSets**,
- provides declarative updates to **Pods**,
- checks on **pod** health
- restarts terminated **pods**.

Deployments should be used instead of directly using **ReplicaSets** unless custom update orchestration is not required or updates themselves are not required.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

Deployment Example (2/2)

<https://kubernetes.io/docs/tutorials/hello-minikube/>

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

Create one container and name it 'nginx' using the field, `.spec.template.spec.containers[0].name`

Deployment named nginx-deployment is created

The Deployment creates three replicated Pods

Defines how the Deployment finds which Pods to manage

Pods run one container, **nginx**, on the nginx Docker Hub image, version 1.14.2

ReplicaSet

<https://kubernetes.io/docs/concepts/workloads/controllers/replicaset/>

The purpose of a **ReplicaSet** is to maintain a stable set of replica **Pods** running at all times. It is often used to guarantee the availability of a specified number of identical **Pods**.

A **ReplicaSet** will dynamically drive the **cluster** back to the predetermined desired state via creation of new **Pods** to keep an application running.

Use **kubectl apply -f <URL>** to apply a template.

Pods can also be added without a template.

A ReplicaSet is defined with fields, including:

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: frontend
  labels:
    app: guestbook
    tier: frontend
spec:
  # modify replicas according to your case
  replicas: 3
  selector:
    matchLabels:
      tier: frontend
  template:
    metadata:
      labels:
        tier: frontend
    spec:
      containers:
        - name: php-redis
          image: gcr.io/google_samples/gb-frontend:v3
```

replicas indicate how many **Pods** should be maintained.

selector specifies how to identify **Pods** it can acquire.

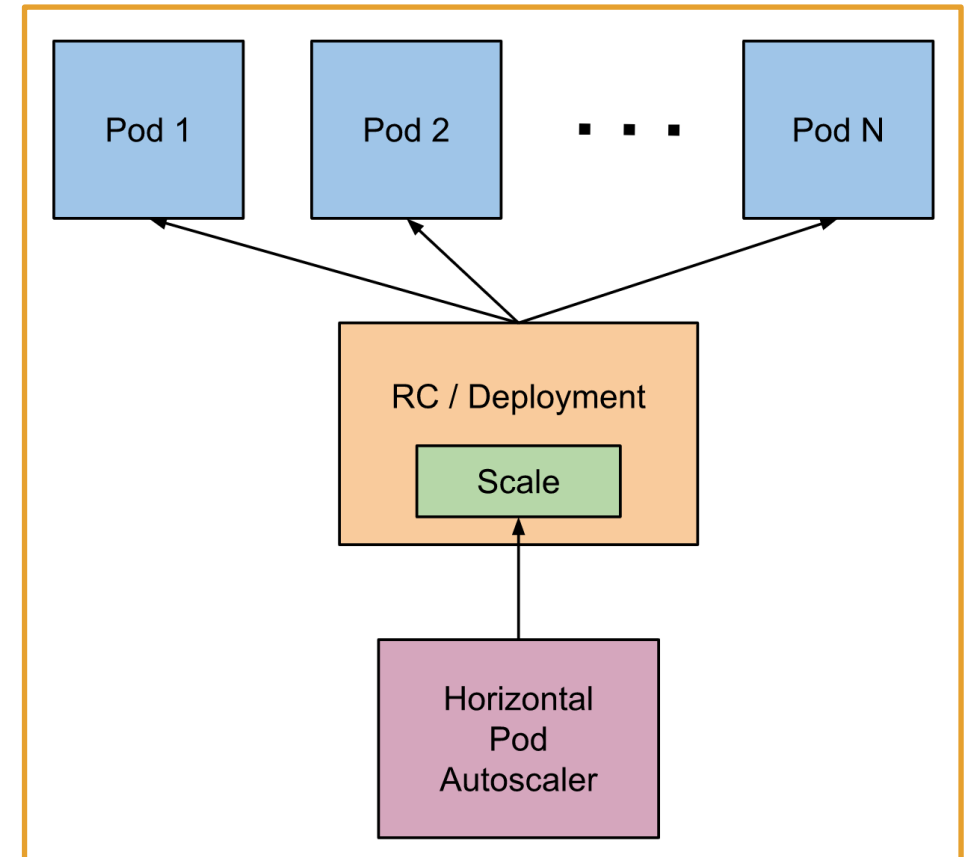
template specifies the data of the new **Pods** that should be created.

AutoScaling

<https://kubernetes.io/docs/tasks/run-application/horizontal-pod-autoscale/>

The ***Horizontal Pod Autoscaler*** is an API resource in the Kubernetes autoscaling API group which automatically scales the number of ***Pods*** in a ***replication controller, deployment, replica set, or stateful set*** based on CPU utilization.

The ***Horizontal Pod Autoscaler*** is implemented as a ***Kubernetes API*** resource and a ***controller*** in a control loop, with a ***period*** maintained by the ***controller manager***. The ***Kubernetes API*** resource determines the behavior of the ***controller***.

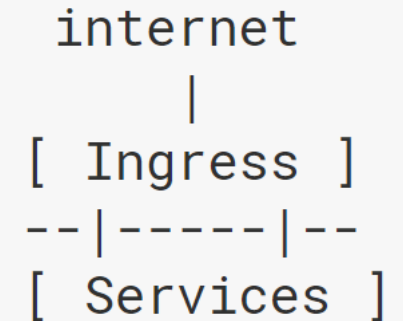


Ingress (1 / 2)

<https://kubernetes.io/docs/concepts/services-networking/ingress/>

In Kubernetes, '*Ingress*' exposes HTTP routes from outside the *cluster* to [services](#) within the *cluster*. It can give *Services* externally-reachable URLs, load balance traffic, terminate SSL/TLS, and offers name-based virtual hosting.

Traffic routing is controlled by rules defined on the *Ingress resource*. An *Ingress Resource* is (usually) a *.yaml* file defining the rules for accessing structures in a *cluster*.



```
apiVersion: networking.k8s.io/v1beta1
kind: Ingress
metadata:
  name: test-ingress
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  rules:
  - http:
      paths:
      - path: /testpath
        pathType: Prefix
        backend:
          serviceName: test
          servicePort: 80
```

Ingress (2/2)

<https://kubernetes.io/docs/concepts/services-networking/ingress/>

As with all other Kubernetes resources, an **Ingress** needs **apiVersion**, **kind**, and **metadata** fields.

Each HTTP **rule** contains:

1. An optional **host**. If no **host** is specified, the rule applies to all inbound HTTP traffic through the IP address specified.
2. A list of **paths**.
3. A **backend** defines a **serviceName** and **servicePort** for each path. It is a combination of **Service** and **port** names.

```
apiVersion: networking.k8s.io/v1beta1
kind: Ingress
metadata:
  name: test-ingress
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  rules:
    - http:
        paths:
          - path: /testpath
            pathType: Prefix
            backend:
              serviceName: test
              servicePort: 80
```

The name of an Ingress object must be a valid DNS subdomain name

Ingress frequently uses annotations to configure some options depending on the Ingress controller

The Ingress **spec** has all the information needed to configure a load balancer or proxy server.


MiniKube

<https://kubernetes.io/docs/setup/learning-environment/minikube/>
<https://kubernetes.io/docs/tasks/tools/install-minikube/>
<https://kubernetes.io/docs/tasks/tools/#minikube>

The ***Minikube CLI*** provides basic bootstrapping operations for working with your cluster, including ***start***, ***stop***, ***status***, and ***delete***.

Minikube:

- is a tool that lets you run Kubernetes locally.
- runs a single-node ***Kubernetes cluster***
- runs on your personal machine (Windows, macOS, Linux PC).
- is great for trying out Kubernetes, or for daily development work.



Hello-Node Tutorial Step-By-Step(1 / 2)

<https://kubernetes.io/docs/tutorials/hello-minikube/>

- Create a Deployment that manages a Pod which will run a container based on the provided Docker Image with:
 - `kubectl create deployment hello-node --image=k8s.gcr.io/echoserver:1.4`
- See the deployment with:
 - `kubectl get deployments.`
- See the Pod with:
 - `kubectl get pods.`
- See cluster events with a:
 - `kubectl get events.`
- See the kubectl configuration with:
 - `kubectl config view.`
- Expose the **Pod** as a Kubernetes **Service** to make it visible from outside the **Cluster** with `type=LoadBalancer` as the expose keyword.
 - `kubectl expose deployment hello-node --type=LoadBalancer --port=8080`

Hello-Node Tutorial Step-By-Step(2/2)

<https://kubernetes.io/docs/tutorials/hello-minikube/>

- View the service you just created with:
 - `kubectl get services`.
- External cloud providers get an external IP to access the service.
- Select + → Select Port to View on Host 1 → Enter the 5 digit port # after the :
- Take a look at available Add-Ons with:
 - `minikube addons list`.
- Enable the metrics-server add-on with:
 - `minikube addons enable metrics-server`.
- View the Pod with the service you just created with:
 - `kubectl get pod, svc -n kube-system`.

Hello-Node Tutorial Step-By-Step(3 / 3)

<https://kubernetes.io/docs/tutorials/hello-minikube/>

- Disable the metrics-server with:
 - `minikube addons disable metrics-server.`
- Delete the service with:
 - `kubectl delete service hello-node`
- Delete the deployment with:
 - `kubectl delete deployment hello-node`
- Stop Minikube with:
 - `minikube stop`
- Delete the Minikube VM with:
 - `minikube delete`

More Tutorials

<https://kubernetes.io/docs/tutorials/>

- <https://docs.microsoft.com/en-us/azure/aks/tutorial-kubernetes-prepare-app>
- <https://kubernetes.io/docs/tutorials/stateful-application/mysql-wordpress-persistent-volume/>
- <https://www.digitalocean.com/community/curriculums/kubernetes-for-full-stack-developers>
- <https://cloud.google.com/kubernetes-engine/kubernetes-comic/>

Assignment (1 / 2)

<https://kubernetes.io/docs/tutorials/kubernetes-basics/>

- Create 8 quiz questions along with four plausible answers each.
- These 8 questions will coincide with the 6 chapters of the above tutorial and the 2 lecture pdf's.
- That means one question will be from each chapter and 1 from each pdf.
- Turn these questions in to your original (first) Batch GitHub repo in a file found at the root of your repo titled '**KubernetesQuiz_FnameLname.txt**'
- These questions will be included in a quiz that you will be given tomorrow and will serve as a primer for the next QC.
- Make sure to include a full range of questions from definitions of elements of Kubernetes structural questions to process to data flow to specific commands in the CTL, etc.
- Also include a link to the docs page where the answer to your question is found.

Assignment (2/2)

<https://kubernetes.io/docs/tutorials/kubernetes-basics/>

Questions format example.

2. This is a DB that stores info about all nodes and clusters in Key:Value pairs.

<https://kubernetes.io/docs/concepts/overview/components/#etcd>

- a. ETCD Cluster (correct)
- b. Worker nodes
- c. Kube-ApiServer
- d. Kuberlet